IBM **TechXchange**

2025 Pre-conference watsonx Hackathon

IBM

IBM TechXchange 2025 Pre-conference watsonx Hackathon

AI agents are intelligent systems that can perceive their environment, reason about it, and take actions to achieve specific goals. These agents are capable of operating autonomously or collaboratively, making them ideal for solving complex real-world problems.

In this hackathon, participants are invited to explore the power of agentic AI using IBM's enterprise ready platforms **IBM watsonx.ai** and **IBM watsonx Orchestrate**. These tools provide everything you need to build, deploy, and manage intelligent AI agents that can plan, adapt, and act in dynamic environments.

IBM watsonx.ai is a powerful AI studio that supports the development of agentic AI solutions using a wide range of foundation models. Through the Prompt Lab, participants can experiment with IBM's Granite models and other leading models to create agents that understand and respond to natural language. IBM watsonx.ai can also serve as an inference provider for your agents, allowing them to generate responses, make decisions, and interact with users or systems intelligently.

IBM watsonx Orchestrate is a no code and low code platform designed to orchestrate AI agents across business workflows. With watsonx Orchestrate, you can create intelligent workflows where agents think, act, and adapt to deliver real business value.

Contents

The hackathon expectation 4
Note on data sets before you begin 4

Get started with IBM watsonx 5
Access your IBM Cloud account 5
Access IBM watsonx 7

1. IBM watsonx.ai 7
Note on IBM watsonx.ai service usage 7
Note on available watsonx.ai services 7
Access Prompt Lab on watsonx.ai 8
Work with the watsonx.ai Prompt Lab 10
Prompt Lab editor 10
Selecting an AI model 10
Programmatic access (API/SDK) 11
watsonx.ai AI agent libraries and tutorials 16
Quick start hands-on exercises 16
Save your Prompt Lab session 17
Save your work on watsonx.ai 18

2. IBM watsonx Orchestrate 19
Note on available watsonx Orchestrate services 19
Access watsonx Orchestrate 19
Discover the catalog 20
Create agents 20
Expanding the abilities of your agent 20
Using agents in the chat 21
Managing connections 21
Building tools 21
Quick start hands-on exercise 21

3. IBM watsonx Code Assistant (OPTIONAL) 21
Note on available watsonx Code Assistant services 21
Access watsonx Code Assistant (WCA) 22
Installing the WCA IDE extension 24
Using the WCA IDE extension 25

Appendix 25

Example use cases 25
Build with agentic AI challenge 25
Sustainable cities and communities challenge, brought to you by Call for Code 26

watsonx Orchestrate hands-on exercise 27
Overview 27
What you'll need 27

Part 1: Build and engage with a new AskBenefits agent 27
Step 1. Create a new AskBenefits agent 27
Step 2. Import and attach the tools 28
Step 3. Upload knowledge document 31
Step 4. Test the agent 33
Step 5. Testing the agent 34
Step 6. Identify the knowledge gap 34

Part 2: Create a new collaborator agent – AskDental 35
Step 1. Create a new AskDental agent 35
Step 2. **Add AskDental agent as a collaborator** 38
Step 3. **Testing scenarios** 40

The hackathon expectation

In this hackathon, participants are expected to design and build a **proof-of-concept agentic AI solution** using **IBM watsonx.ai, IBM watsonx Orchestrate, or both**. Your solution should align with one of the following themes:

1. **Build with agentic AI challenge**
   Identify a real-world everyday problem and develop a working proof-of-concept that uses agentic AI to solve it. Focus on high impact tasks or workflows where AI agents can demonstrate autonomy, collaboration, and adaptability.

2. **Sustainable cities and communities challenge presented by Call for Code**
   Create an innovative agentic AI proof-of-concept that addresses a challenge under United Nations Sustainable Development Goal (SDG) 11 which focuses on sustainable cities and communities. Think about how AI agents can improve urban living whether through smarter infrastructure, efficient resource management, or inclusive community services.

You are required to use **IBM watsonx.ai, IBM watsonx Orchestrate, or both**. These platforms provide the foundation for building and orchestrating intelligent AI agents.

Participants may also **optionally** use **IBM watsonx Code Assistant** to accelerate development. Generate code from natural language, integrate APIs, and streamline backend logic to support your agentic AI solution.

Example use cases

Refer to example use cases to help you get started with your solution idea.

Note on data sets before you begin

Participants are required to bring their own datasets to build the solution aligning to your use case. As you collect data for your project, you'll want to use the best practices. Here are some helpful tips:

- Teams are responsible for ensuring data is compliant.
- Data from public websites may be used, if the terms allow for commercial use, but please keep a list of the websites you use.
- Do not use data or assets containing company confidential data, or any other data without permission from the data owner. Teams are responsible for getting approval.
- Do not use any client data.
- Do not use any data containing personal information (PI).
- Do not use data obtained from social media.

Get started with IBM watsonx

To access and use the watsonx products for this hackathon, participants must be registered for the hackathon and have access to the hackathon site. Once you have access to the hackathon site, follow the instructions on the "**Complete the hackathon**" page to request a pre-configured hackathon IBM Cloud account for your team.

Access your IBM Cloud account

Once your team has been provisioned a hackathon IBM Cloud account, all team members will receive an email invite to join the cloud account. Follow the steps below to access your team's cloud account:

1. Check the email inbox you used to register for the hackathon and open the email you received from the IBM Cloud team about joining your cloud account. Please check your junk/spam folders if you are not able to find the email in your inbox. You can also quickly search for "IBM Cloud" to locate the email.

2. Click the **Join Now** button seen in that email. A new browser tab will open with the cloud account sign up page.

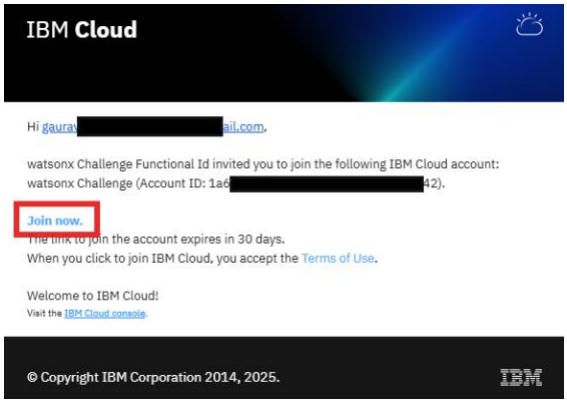

3. Review your account and personal information. Read and accept the Account notice and click the **Join Account** button.

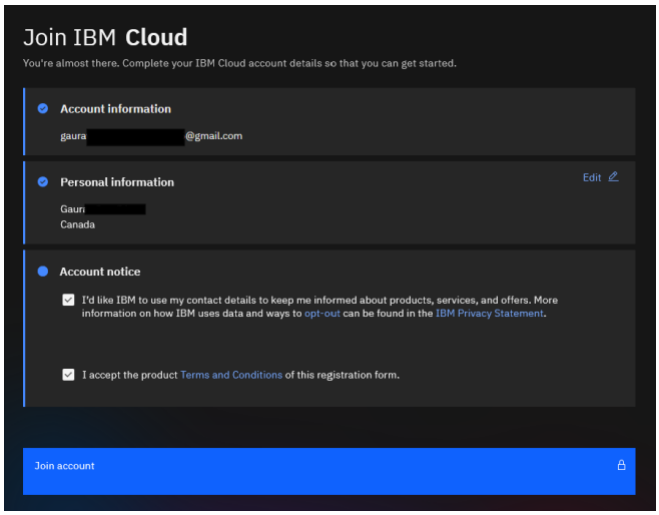

4. Complete the authentication process by clicking the Continue button.

5. After you authenticate successfully, you will be taken to the IBM Cloud dashboard.

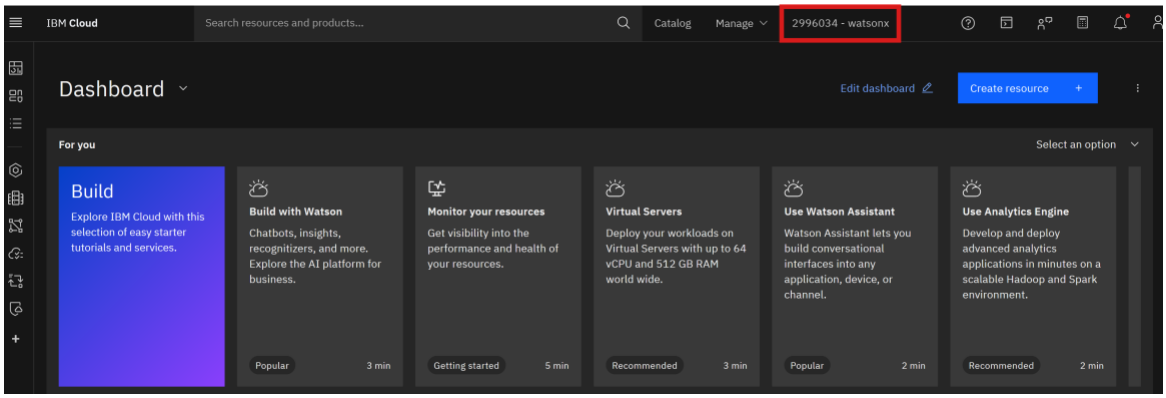

6. If you have an existing personal IBM Cloud account for the same email/IBMid, sometimes you will be directed to your personal account. In this case, please switch your account to the **xxxxxxx - watsonx** account. Select your account drop-down at top-right of the dashboard and select watsonx account. Refer to the below image on switching accounts in your cloud dashboard.

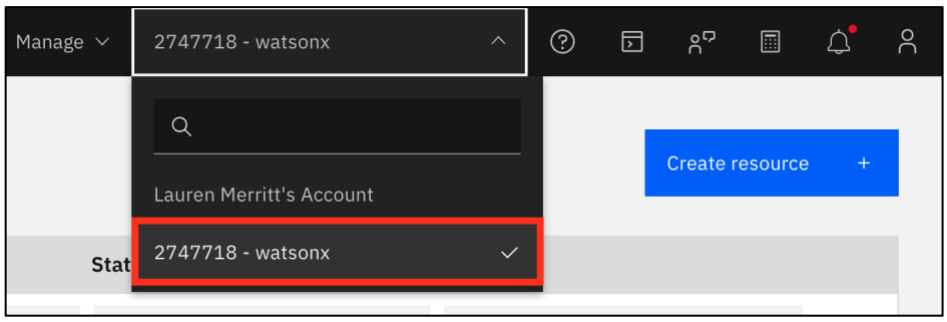

Access IBM watsonx

To begin building your agentic AI solution, explore the capabilities and resources for each IBM watsonx product.

1. watsonx.ai
2. watsonx Orchestrate
3. watsonx Code Assistant (Optional)

1. IBM watsonx.ai

After successfully joining the IBM Cloud account, you can now access watsonx.ai to work on the platform and build your solution.

Note on IBM watsonx.ai service usage

For this hackathon, **$100 credits** will be automatically applied on the provisioned **IBM watsonx.ai platform**. This should be sufficient for designing and creating compelling submissions.

You will receive periodic email notifications about your **credit consumption** at the following usage levels: **25%, 50%,** and **80%**. Once you reach **100% usage**, your account will be **suspended**. You can appeal the suspension by completing the form shared in the account suspension notification email.

Please note that these email notifications are sent **once per hour**, so there is a possibility that you may **exhaust all your credits before receiving an alert**.

Please plan to use the watsonx.ai efficiently and back up your work accordingly. Refer **tips to work efficiently on watsonx.ai platform (Tokens and CUH explained)** and **saving your work**.

Important:

- **Foundation model inferencing** consumes tokens, which are measured as Resource Units (RUs). **1,000 tokens = 1 RU,** and each RU costs **$0.0001 USD.**
  Learn more about tokens and tokenization.
- If you are using **Jupyter Notebook editor on watsonx.ai,** consider selecting **a lower runtime environment** to avoid high resource consumption and quickly depleting your credits. Notebook runtimes are billed based on **Capacity Unit Hours (CUH)** at a rate of **$1.02 USD per CUH.**
  Learn more about capacity unit hours and watsonx.ai Studio pricing plans.

Note on available watsonx.ai services

The watsonx.ai platform is **pre-configured with only the services required** to complete the hackathon. If you notice a permission/access issue for any service or the cloud catalog, then they are not required/available for this hackathon.

These features/capabilities are out of scope for this hackathon:

- Agent Studio (Beta)
- Deploy on IBM Cloud/watsonx.ai (including Deployment space)
- Bring your own model
- Fine tuning models
- AutoAI pipeline
- SPSS Modeler
- Federated Learning
- Cloud Object Storage service

DO NOT USE the below listed models as they are out of scope for the hackathon and can negatively impact the judgment of your project submission.

- llama-3-405b-instruct
- mistral-large
- mixtral-8x7b-instruct-v01
- mistral-medium-2502
- mistral-small-3-1-24b-instruct-2503
- pixtral-12b

The hackathon provisioned IBM Cloud account will be deactivated after the completion of the hackathon. Please plan to save your work at the end of the hackathon.

Access Prompt Lab on watsonx.ai

After successfully joining the IBM Cloud account, you can now access the Prompt Lab on watsonx.ai platform to work with the AI models supported on the platform and build your solution.

1. Log in to the watsonx.ai platform (https://dataplatform.cloud.ibm.com/wx/home?context=wx) with the email you used to access your IBM Cloud account.
2. After successful authentication, you will see “Welcome to watsonx”. You can either take the tour or skip it.
3. Next, you will see the watsonx.ai dashboard. Ensure the name of the account is “**xxxxxxx – watsonx**” and the region is “**Dallas**”.

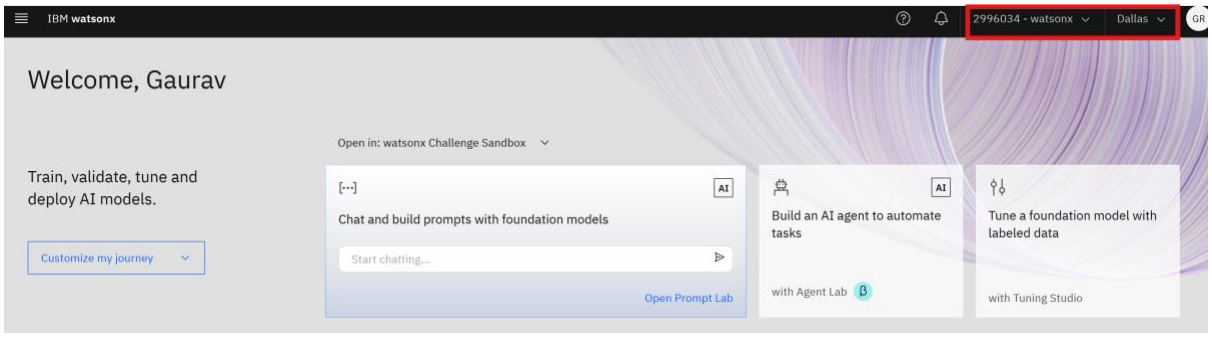

4. Select the "**Open Prompt Lab**" button on the "Chat and build prompts with foundational models" widget.

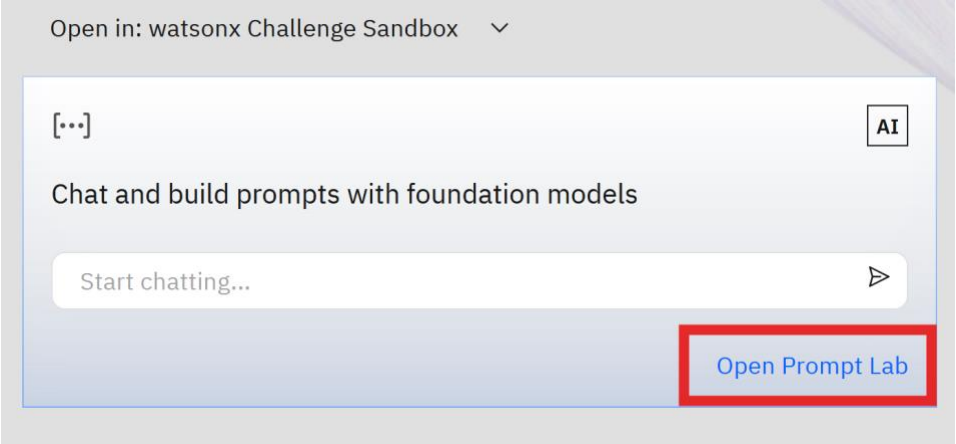

5. The "Welcome to Prompt Lab" tour will be displayed. You can take the tour to get a quick introduction or skip it.
6. The Prompt Lab Editor opens with a chat window to get you started with the prompt session.

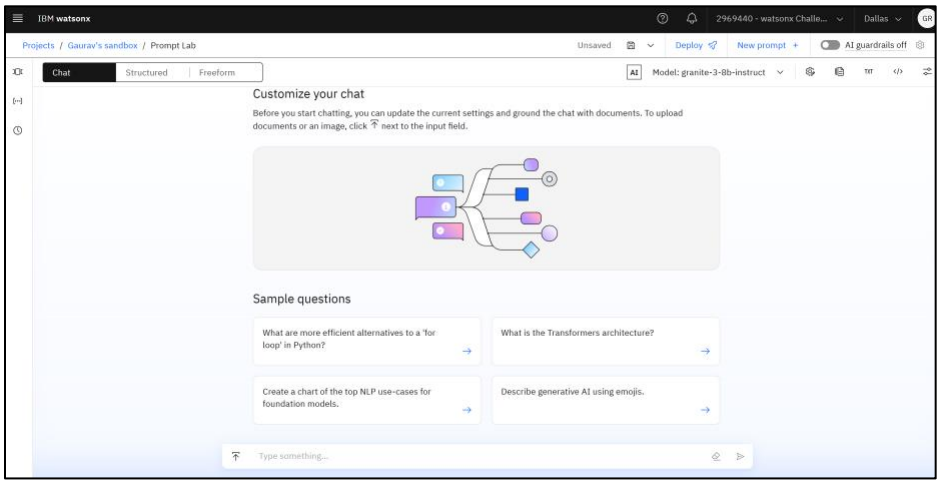

Work with the watsonx.ai Prompt Lab

The watsonx.ai Prompt Lab is an easy-to-use prompt engineering interface where you can experiment prompting different AI foundation models, explore sample prompts, tune model parameters, integrate applications with an API endpoint, and save and share your best prompts.

Take a tour of the Prompt Lab and try the interactive demo.

You can access and use the AI models to build your innovative solution using Prompt Lab.

Prompt Lab editor

In the Prompt Lab, you can experiment with prompting different foundation models, explore sample prompts, as well as save and share your best prompts. The Prompt Lab editor is a great place to experiment and iterate with your prompts. Try the quick start lab.

However, you can also prompt foundation models in watsonx.ai programmatically. Refer to "Programmatic access (API/SDK)" section.

Selecting an AI model

A default AI model will be pre-selected in the Prompt Lab editor. You can either use the same model or change to a different model. To select a different model:

1. Select the AI Model drop-down menu at the top-right of the editor and select **View all foundation models**.

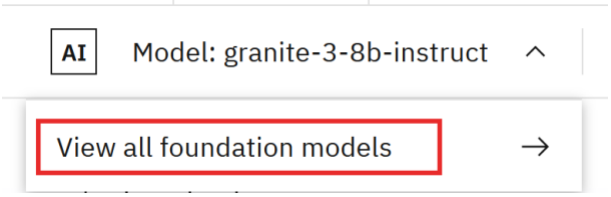

2. The **Select a foundation model** widget will appear. Clear the filters to see all the available models. You can use the filters to choose the right model for your solution building. You can select a model tile to learn about the model and use it. If you are limited to only "Chat" supported models, change the Prompt Lab editor to **Structured** or **Freeform** view and try selecting the models to see all the available model options.

   Important: DO NOT USE the below listed models as they are out of scope for the hackathon and can negatively impact the judgment of your project submission.

   - llama-3-405b-instruct
   - mistral-large
   - mistral-medium-2502

- mixtral-8x7b-instruct-v01
- mistral-small-3-1-24b-instruct-2503
- pixtral-12b

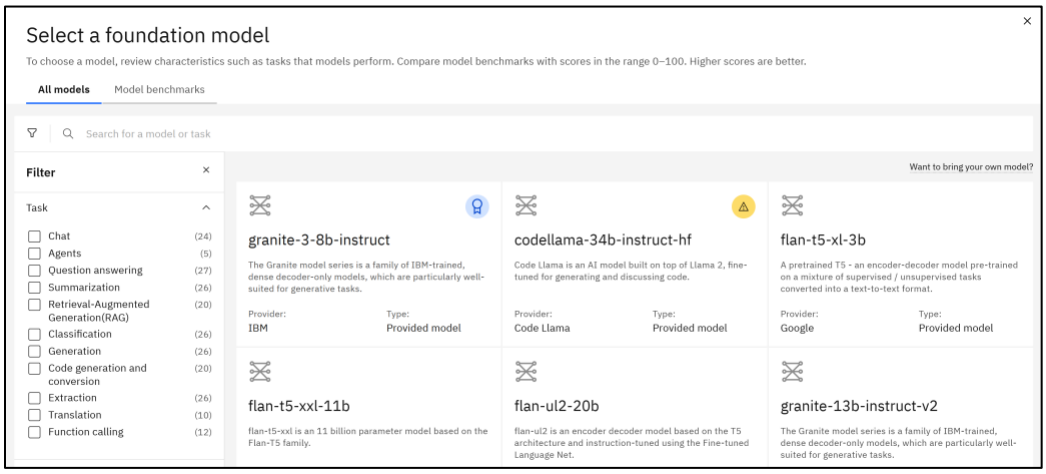

To understand how models can address your use case, including information on model modalities, supported languages, tuning, and indemnification, see our product documentation on choosing a model.

Note: Bigger models are not always better. Learn why smaller models can be better and more cost effective.

Programmatic access (API/SDK)

You can inference the watsonx.ai models with API or SDK requests.

Developer access information

To use the supported watsonx.ai APIs/SDKs, you will need three values: a **project ID**, an **endpoint URL** and an **API key**.

- Go to watsonx.ai home page.
- Select the hamburger icon in the top left of the screen.

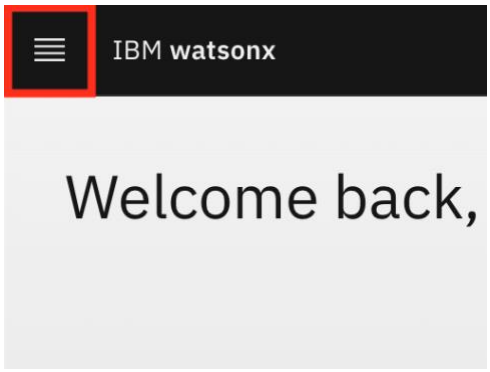

- From the drop down select *Support* > *Developer Access*.

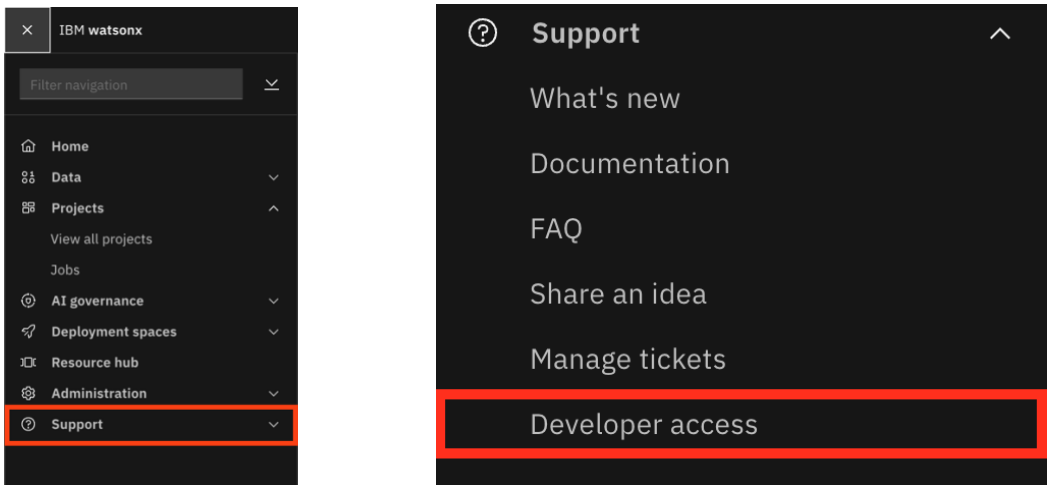

- Once on the Developer access page, in the Base URL section a default **watsonx.ai endpoint URL** will be displayed for the Dallas region. Ensure the region is always set to **Dallas** at the top right of the page.

- In the "**Project or space ID**" section you will also find your **Project ID**. Select from the "**Project or space**" drop-down, the "**watsonx Challenge Sandbox**" option if it is not already selected. A **project ID** will be displayed.
  Note: A space ID is **not required** as it is out of scope for the hackathon.

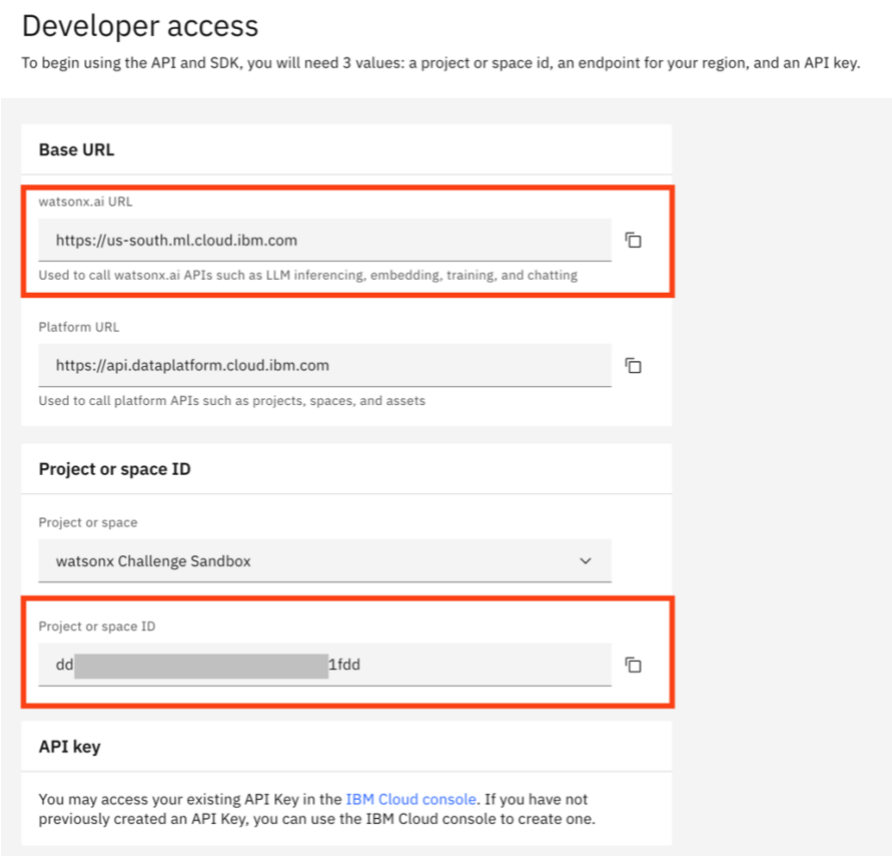

- In the "**API key**" section you will find instructions on how to access and/or create your API key. Navigate to the IBM Cloud console page linked in this section.

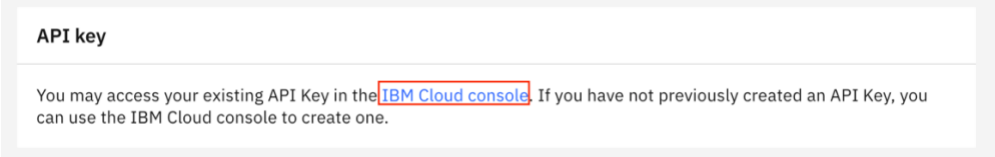

- If you do not already have an existing API key you would like to use, select the blue "**Create**" button.

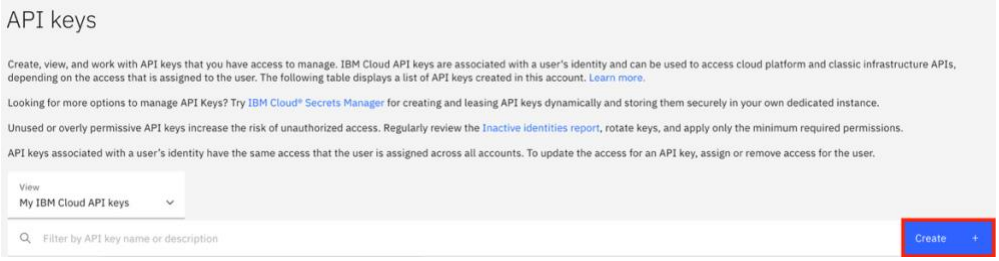

- A **Create API key** dialog box will be displayed. Enter a name, provide optional description and choose the "Disable the leaked key" option. Click the "**Create**" button.

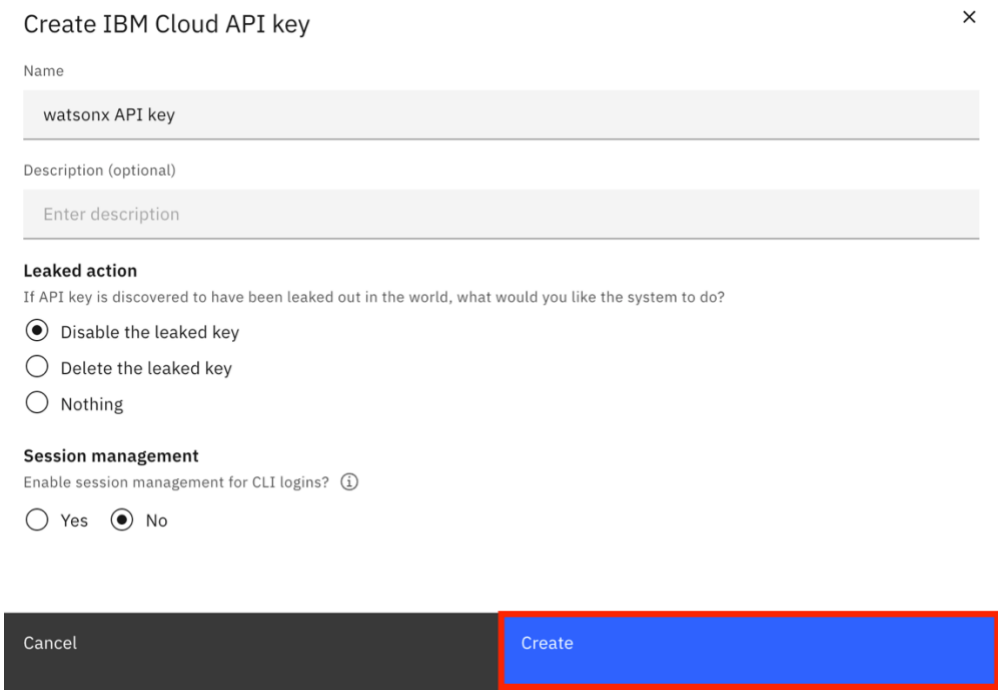

- An API key will be created successfully. Copy the API key and save it safely to use for calling the API/SDK. You can also download and save the file in a secure path in your system.

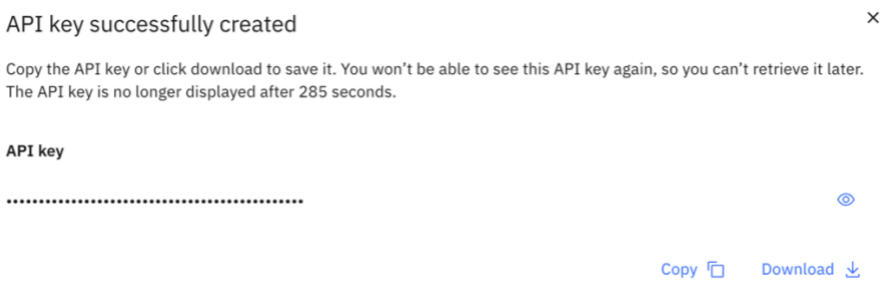

watsonx.ai programmatic options

There are multiple options to help you get started using watsonx.ai APIs/SDKs.

Option 1: Prompt Code on Prompt Lab

Refer to the access prompt code instructions to learn how to quickly get access to the text generation API within the watsonx.ai Prompt Lab.

Option 2: Different watsonx.ai API capabilities

Explore and leverage different watsonx.ai API capabilities in your solution.

- Chat
- Agent-driven chat
- Tool calling
- Text generation
- Time series
- Text rerank
- Embeddings
- Text extraction

Refer supported API functionality by model here.

Access the prompt code (API) from Prompt Lab editor

To prompt an AI model programmatically, you can view and copy the prompt code by selecting the **View code** icon </> at the top-right of the prompt lab editor.

The prompt code is available as a Curl, Node.js and Python.

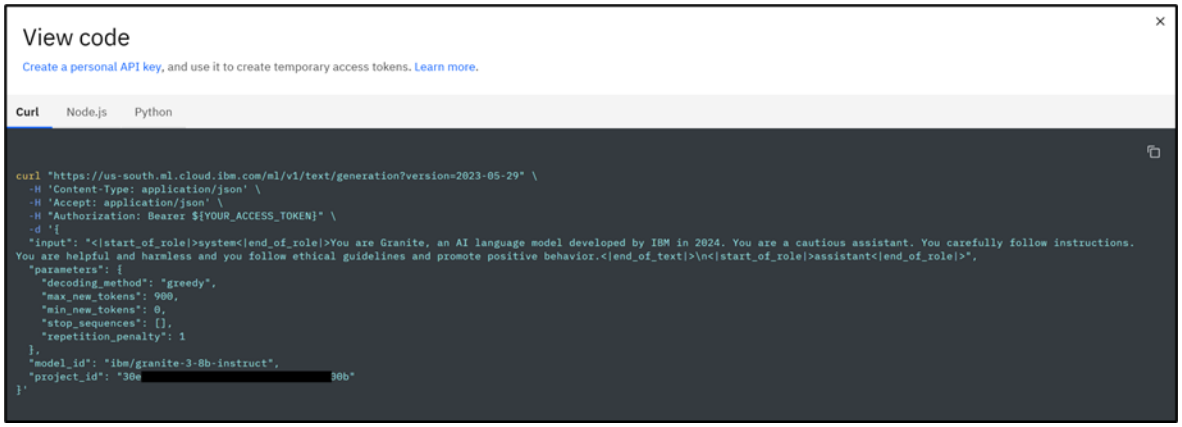

You will require an IAM access token to authorize the prompt code and need to replace **${YOUR_ACCESS_TOKEN}** placeholder in the prompt code. You can create an IAM access token using an API key.

- **API key:**

  Refer to Developer access information to get an API key.

- **Generate IAM Access Token:**

  Programmatically generate an IAM access token with the API key using the following cURL command:

```
curl -X POST 'https://iam.cloud.ibm.com/identity/token' -H 'Content-Type: application/x-www-form-urlencoded' -d 'grant_type=urn:ibm:params:oauth:grant-type:apikey&apikey=MY_APIKEY'
```

- **curl -X POST** → Specifies an HTTP **POST** request.
- **URL ("https://iam.cloud.ibm.com/identity/token")** → The endpoint to request an authentication token from IBM Cloud.
- **-H "Content-Type: application/x-www-form-urlencoded"** → Sets the request header to indicate that the data is sent in form-encoded format.
- **-d (Data Payload)** → Sends the required data:
- **grant_type=urn:ibm:params:oauth:grant-type:apikey** → Specifies the OAuth grant type as API Key.
- **apikey=MY_IBM_CLOUD_API_KEY** → Replace MY_IBM_CLOUD_API_KEY with your actual IBM Cloud API key.

```
Expected Response:

{
    "access_token": "eyJhbGciOiJIUz......sgrKIi8hdFs",
    "refresh_token": "not_supported",
    "token_type": "Bearer",
    "expires_in": 3600,
    "expiration": 1473188353,
    "scope": “ibm openid”
}
```

> **Note**: An IAM token is valid for up to 60 minutes, and it is subject to change. When a token expires, you must generate a new one. Use the property "*expires_in*" for the expiration of the IAM token that you have just created.

watsonx.ai AI agent libraries and tutorials

Explore the watsonx.ai supported AI agent framework libraries and tutorials to help you get started building your AI agent solution.

- LangChain
- LangGraph
- LlamaIndex
- CrewAI
- BeeAI
- Python SDK
- Node.js SDK

Quick start hands-on exercises

Try the quick start exercises and notebooks for sample use cases to get started with using watsonx.ai.

Important notes:

- Refer to developer access information section to use watsonx.ai credentials as you try the exercises.
- Some of the exercises could include the usage of old model version. You can replace them with newer versions for better performance and output. To check the latest supported Granite models on watsonx.ai, either follow selecting an IBM Granite model on Prompt Lab or refer to supported foundation models on watsonx.ai.
- The hackathon provisioned cloud accounts **do not support solution deployment.** You can run your solution deployment locally on your machine and showcase them in your submissions.
- **Foundation model inferencing** consumes tokens, which are measured as Resource Units (RUs). **1,000 tokens = 1 RU,** and each RU costs **$0.0001 USD.**
  Learn more about tokens and tokenization.
- If you are using **Jupyter Notebook editor on watsonx.ai,** consider selecting **a lower runtime environment** to avoid high resource consumption and quickly depleting your credits. Notebook runtimes are billed based on **Capacity Unit Hours (CUH)** at a rate of **$1.02 USD per CUH.**
  Learn more about capacity unit hours and watsonx.ai Studio pricing plans.

watsonx.ai Prompt Lab app templates:

- LangGraph LLM app template with function calling capabilities (base template)
- LlamaIndex Workflow LLM app template with function calling capabilities (base template)
- CrewAI LLM app template with function calling capabilities (base template)
- arXiv Research agent (community template)
- Agentic RAG LangGraph template (community template)

BeeAI Agent Framework:

- BeeAI framework examples

LangChain and LangGraph:

- Create a LangChain AI Agent in Python using watsonx
- Build a RAG agent using LangGraph to answer complex questions
- Build a LangChain agentic RAG system using the Granite model in watsonx.ai
- Use watsonx, and LangChain Agents to perform sequence of actions
- Use watsonx, and LangChain to make a series of calls to a language model
- arXiv Research agent
- Base LangGraph LLM app template with function calling capabilities

LlamaIndex:

- Use watsonx and LlamaIndex for Text-to-SQL task
- Use watsonx, and `llama-3-1-70b-instruct` and LlamaIndex to make simple chat conversation and tool calls
- LlamaIndex Workflow LLM app template with function calling capabilities

CrewAI:

- Leveraging CrewAI and IBM watsonx
- Build an agentic framework with CrewAI memory, i18n, and IBM watsonx.ai
- Base CrewAI LLM app template with function calling capabilities

Save your Prompt Lab session

You can save your Prompt Lab editor session for later use.

1. At the top of the Prompt Lab screen, select the **Save work** dropdown button and then select the **Save as** option.

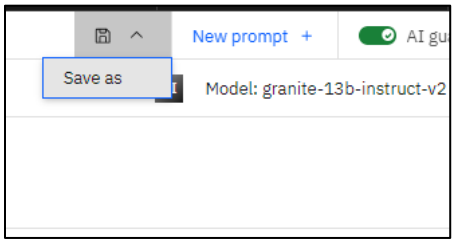

2. A **Save your work** widget will appear. Select **Prompt session** under the **Asset type** option.
3. Enter a **name** and check the **View in project after saving** option under the **Define details** section.
4. Finally, click the **Save** button. Once you save, you will see the saved work under the **Assets** tab

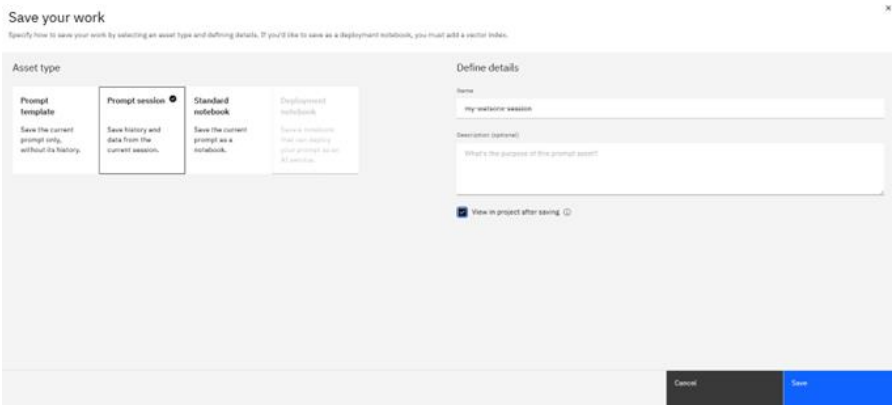

You can also save your work as:

- **Prompt template** to save only the current prompt without its history and selecting a **Task** suitable for your prompting.
- **Notebook** to continue prompting on a Jupyter Notebook environment. Prior knowledge of notebooks and Python programming language would be helpful to work with a Jupyter notebook. Read more about notebooks.

Save your work on watsonx.ai

Make sure to save any work you want to retain for your records. IBM Cloud accounts will be deactivated at the end of the hackathon. Follow the steps below to save your work:

1. Go to your project's 'Overview' tab.
2. Select the 'Export or import project' drop down below the Bell icon in the top menu bar.
3. Click the 'Export project' option. This will open 'Export project to desktop' screen.
4. Select all the assets shown in your project (Work saved as Project session cannot be exported) and click 'Export' on the bottom-right of the screen.
5. The next screen will ask for confirmation that all sensitive information has been removed.
6. Click on 'Continue export'.
7. The download (zip) will be initiated, and the file will be saved on your computer.

2. IBM watsonx Orchestrate

After successfully joining the IBM Cloud account, you can now access the watsonx Orchestrate to work on the platform and build your solution.

Note on available watsonx Orchestrate services

A **Trial** watsonx Orchestrate service is **pre-configured with only the services required** to complete the hackathon. If you notice a permission/access issue for any service or the cloud catalog, then they are not required/available for this hackathon.

These features/capabilities are out of scope for this hackathon:

- Configuring a voice

Access watsonx Orchestrate

1. In your IBM Cloud account dashboard, select the "**Navigation menu**" on the top left of the dashboard and select the "**Resource list**" option.

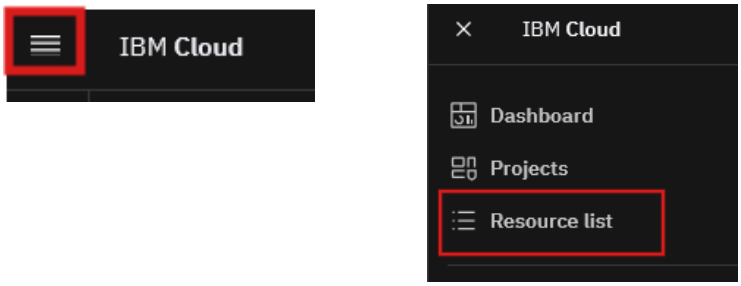

2. Expand the "**AI / Machine Learning**" section and select "**watsonx-Challenge Orchestrate**" service.

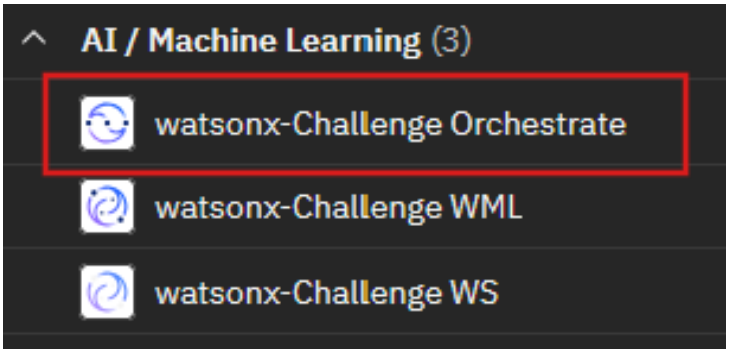

3. You will be navigated to **watsonx-Challenge Orchestrate** service instance dashboard. Select the "**Launch watsonx Orchestrate**" button.

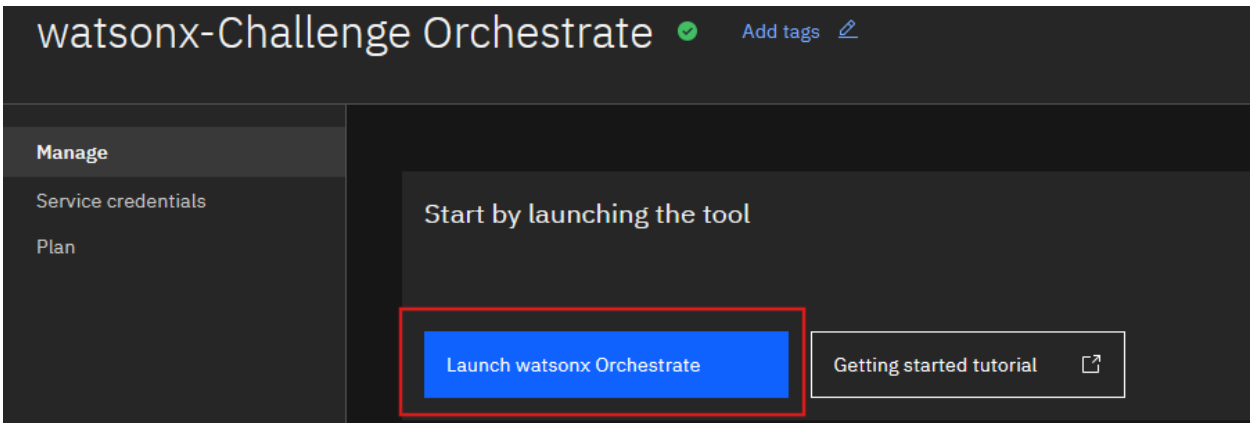

4. You will be navigated to the watsonx Orchestrate platform with a welcome message and a new chat window.

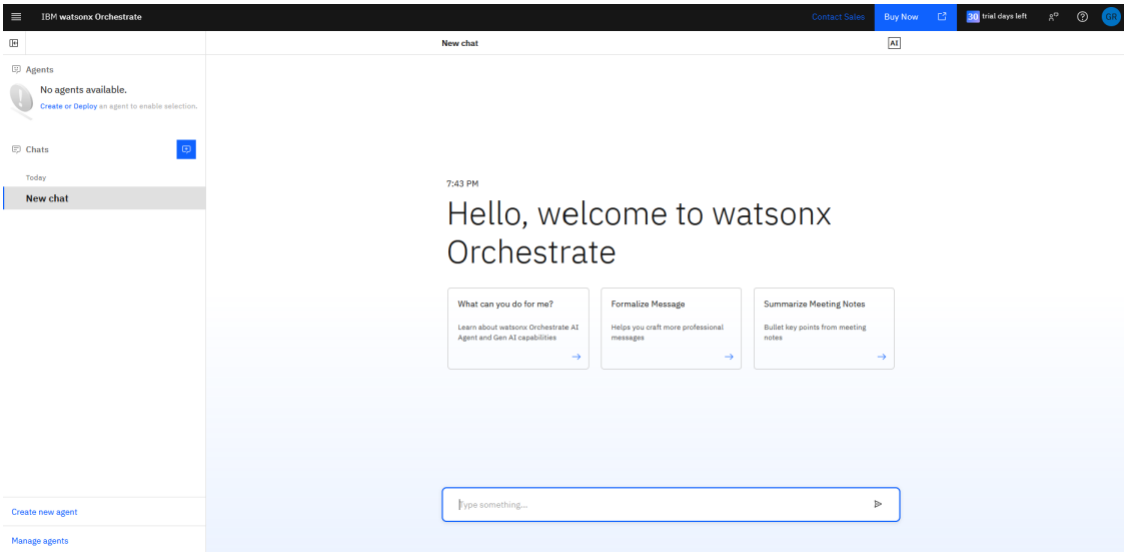

Discover the catalog

The IBM watsonx Orchestrate catalog serves as a vast repository of prebuilt AI agents and tools, tailored to address a wide array of use cases and requirements. This extensive collection helps you to discover agents, tools, or a blend of both that align with your specific needs. Learn more about discovering the catalog here.

Create agents

You can create a new agent in IBM watsonx Orchestrate by either creating from scratch or using a pre-built template. The agent creation process involves defining its name and purpose that is suited to your specific needs. Learn more about creating agents here.

Expanding the abilities of your agent

To expand your agent's abilities, explore the following configuration options:

- Recommendations for agent descriptions

- Customizing agent styles
- Adding knowledge to agents
- Adding agents for orchestration
- Adding tools to agents
- Adding instructions to agents
- Connecting to channels
- Deploying the agent

Using agents in the chat

In IBM watsonx Orchestrate, agents collaborate to automate tasks and manage workflows. Learn more about it here.

Managing connections

To use the external applications within IBM watsonx Orchestrate, you must establish a connection between them which acts as bridge enabling communication between watsonx Orchestrate and the external applications. Learn more about it here.

Building tools

Tools can be used to enhance your agent's capabilities and allow them to perform several operations. You can use pre-built tools, create your own or import existing tools, Learn more about it here.

Quick start hands-on exercise

Try the quick start hands-on exercise to get started with using watsonx Orchestrate.

3. IBM watsonx Code Assistant (OPTIONAL)

After successfully joining the IBM Cloud account, you can now access the watsonx Code Assistant to build your solution.

Note on available watsonx Code Assistant services

A **Trial** watsonx Code Assistant service is **pre-configured with only the services required** to complete the hackathon. If you notice a permission/access issue for any service or the cloud catalog, then they are not required/available for this hackathon.

The Trial WCA service comes with a limit of 150,000 tokens, and additional tokens will not be provided for the hackathon. You can track your service token consumption on WCA service landing page or by selecting the "**Tokens and resource units"** tile on the home page**.**

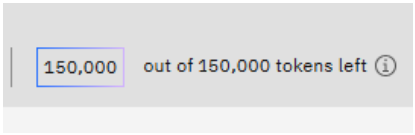

Access watsonx Code Assistant (WCA)

1. In your IBM Cloud account dashboard, select the "**Navigation menu**" on the top left of the dashboard and select the "**Resource list**" option.

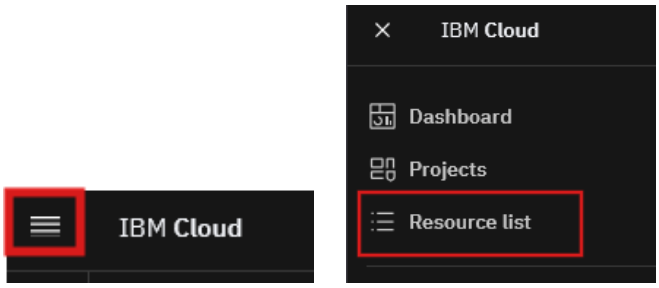

2. Expand the "**AI / Machine Learning**" section and select "**watsonx-Challenge Code Assistant**" service.

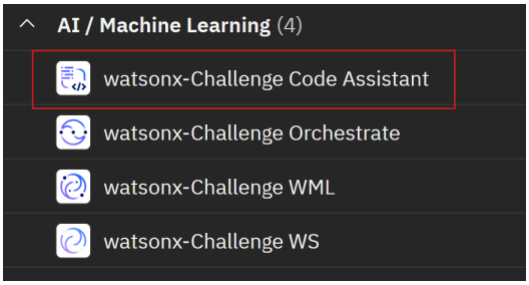

3. You will be navigated to **watsonx-Challenge Code Assistant** service instance dashboard. Select the "**Launch watsonx Code Assistant**" button.

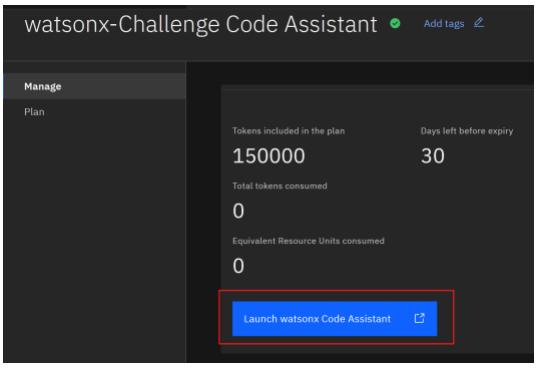

4. You will be navigated to the watsonx Code Assistant dashboard, and a setup widget will be displayed. Select the "**Set up**" button.

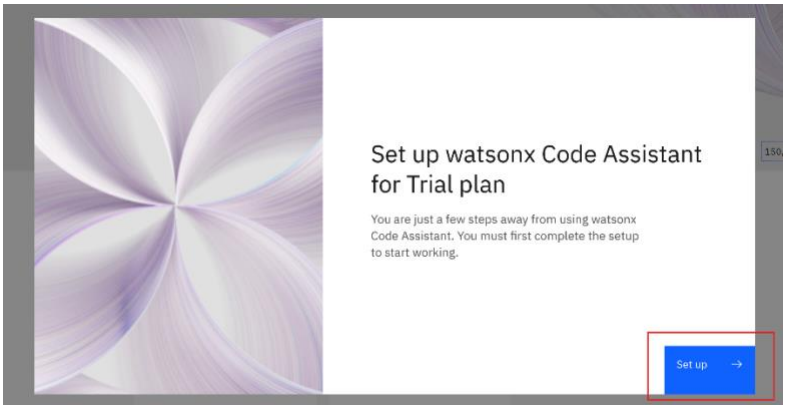

5. Once the setup is completed, you will see "**You're all set**" confirmation. Select the "**Go to home page"** button.

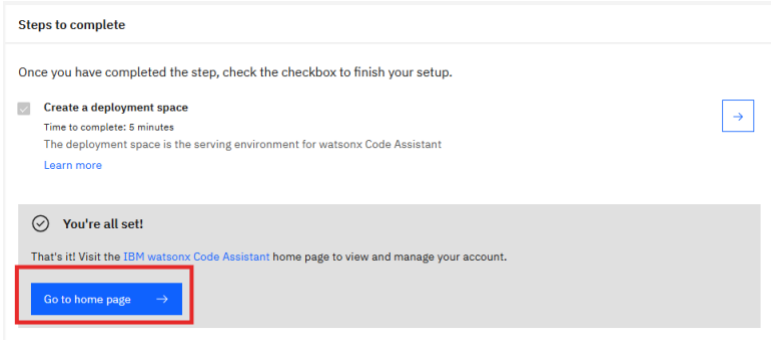

Important: If you do not see the "**You're all set**" confirmation and selection to "Go to home page" button,

i) Click the blue checkbox beside **Create a deployment space**.

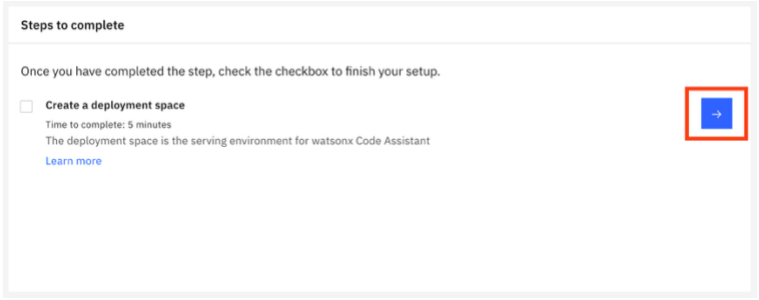

ii) In the "Create a deployment space" dialog box select **Cancel** in the bottom right corner of the page.

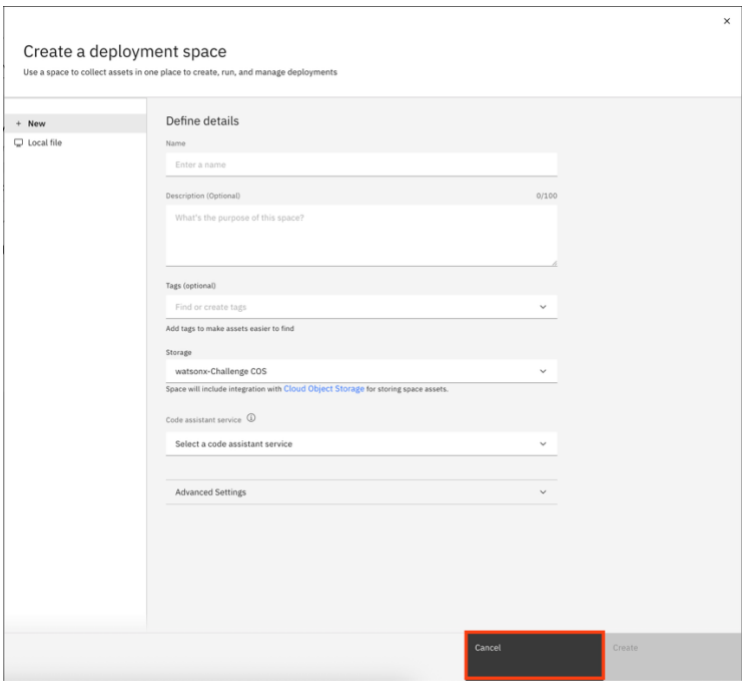

iii) Check the checkbox beside “Create a deployment space”. Once selected you will see the “You’re all set” confirmation. Select the “Go to home page” button.

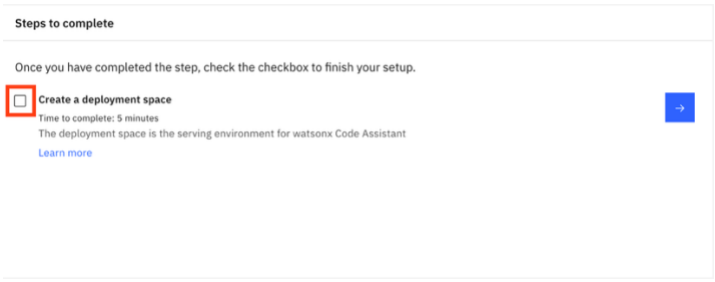

6. You will be navigated to the home page, and a welcome message will be displayed. You can either take the tour or skip it. Now you are all set and ready to install the WCA extension in your preferred IDE.

Installing the WCA IDE extension

The watsonx Code Assistant (WCA) supports extension installation for Visual Studio Code and Eclipse IDEs. Select your preferred IDE to view the installation instructions.

- Visual Studio Code
- Eclipse

Note: An IBM Cloud API key would be required to login and use the extension. Refer to Create an API key instructions to generate one and use it.

Tip: If you decide to use Visual Studio Code, as you complete the log in steps the field to enter your API key can be found in the top-center of the editor screen.

Using the WCA IDE extension

Try the interactive demo.

Refer to the following instructions to learn and try different capabilities of the product:

- Generate code suggestions
- Explaining code
- Documenting code
- Generating unit tests
- Translating code from one language to another
- Prompting guide

Appendix

Example use cases

You are not limited to these ideas, but here are several examples of how you could apply IBM watsonx to solve a specific issue within your chosen theme:

Build with agentic AI challenge

- **AI coding assistant for developer productivity**: An agent that helps software developers plan tasks, generate code snippets, resolve bugs, and learn new APIs by integrating natural language prompts with enterprise codebases and documentation.
- **Code review companion**: An AI agent that reviews pull requests for bugs, style issues, and security vulnerabilities, suggests improvements, and highlights potential regressions before human review.
- **Developer onboarding agent**: An agent that guides new engineers through environment setup, project architecture, and internal tooling by answering natural language questions and automating configuration tasks.
- **Continuous Integration and Continuous Deployment (CI/CD) monitoring and debugging assistant**: An agent that watches CI/CD pipelines, explains build failures in natural language, and suggests or auto-tests fixes by referencing build logs, commit history, and config files.
- **AI debugging and learning assistant**: An agent that explains compiler or runtime errors in simple terms, suggests fixes, and links to relevant documentation or tutorials — all tailored to the learner's skill level.
- **Academic research and study assistant**: An agent that helps university students structure research papers, track sources, create study schedules, and quiz themselves using lecture notes and external materials.
- **AI career copilot for job seekers**: An AI agent that guides individuals through job discovery, resume creation, skill development, and interview preparation.
- **Smart meeting synthesizer and follow-up agent**: An autonomous agent that joins meetings, extracts

key decisions and action items using watsonx.ai, and follows up with stakeholders to track progress.

- **Personalized learning path advisor for technical professionals**: An agent that recommends upskilling paths based on an individual's role, goals, and learning style by integrating company learning and development catalogs with open resources and certifications.
- **Travel and errand planning assistant**: An AI agent that dynamically plans daily errands, appointments, and travel routes by optimizing for traffic, store hours, and personal preferences.
- **AI companion for busy households**: An AI agent that helps parents manage family logistics by coordinating school events, meal planning, to-do lists, and reminders across multiple family members.
- **Adaptive AI tutor for neurodiverse learners**: An intelligent tutor that personalizes lessons and support for students with learning differences such as ADHD or dyslexia.
- **Micro-business operations assistant**: An AI agent that automates key business tasks like invoicing, customer communication, and marketing for solopreneurs.

Sustainable cities and communities challenge, brought to you by Call for Code

- **Urban air quality sentinel**: Develop an agent that monitors urban air quality in real-time and autonomously triggers mitigation strategies such as rerouting traffic, notifying residents, or activating urban greening systems.
- **Housing rights navigator**: Create a multilingual AI agent that assists vulnerable residents in identifying, applying for, and protecting their rights to safe and affordable housing while alerting legal or support services in case of eviction risks.
- **Inclusive urban planning co-pilot**: Build an agent that reviews urban development proposals, detects potentially exclusionary language or impacts, and recommends inclusive changes to promote equitable planning outcomes.
- **Resilient infrastructure watchdog**: Create an infrastructure agent that continuously analyzes data from roads, bridges, and utilities to predict potential failures and notify relevant authorities for proactive maintenance.
- **Green space access optimizer**: Create an agent that evaluates access to public green spaces, identifies underserved areas, and proposes data-driven locations for new parks or upgrades to existing ones.
- **Dignified housing upgrade platform**: Develop a digital platform that enables residents of informal settlements to co-design and fund dignified housing upgrades in partnership with local governments and NGOs.
- **Civic participation dashboard**: Create a civic engagement dashboard that aggregates and analyzes public feedback, ensuring inclusive participation in urban planning processes and supporting transparent governance.
- **Mobility equity analytics engine**: Build a data analytics platform that maps gaps in public transport access and recommends affordable mobility solutions like micro-transit or bike-share in underserved areas.
- **AI-based tenant assistant for renters**: An agent that helps renters manage lease agreements, submit maintenance requests, track rent payments, and communicate with landlords or property managers.
- **Urban sustainability impact tracker**: Launch a digital certification and impact tracking platform that helps buildings and neighborhoods monitor sustainability metrics and benchmark progress toward equitable, green living standards.

watsonx Orchestrate hands-on exercise

Overview

In this activity, you'll step into the role of an AI solution developer working within the IBM watsonx Orchestrate platform. The lab is designed to simulate a real-world healthcare use case by guiding you through the process of creating, configuring, and deploying intelligent agents from scratch. This activity demonstrates how intelligent agents can be configured and deployed in watsonx Orchestrate to address healthcare-specific use cases. Upon the completion you will:

- Understand the end-to-end agent creation process
- Learn how agentic AI uses context, memory, and tools to enable adaptive, natural multi-turn interactions
- Be familiar with OpenAPI-based tool integration
- Get practical experience with multi-agent solutions
- Gain experience testing and validating agent behaviour in practical scenarios

You'll be completing two core activity parts:

1. **Build your own AskBenefits agent** – a digital agent that answers questions about medical procedures and health plan benefits. It detects if a procedure is past, future, or overdue and uses the right tools to respond. It can retrieve past procedure data, flag overdue care, guide upcoming procedures, help schedule/reschedule appointments and showcase HR policies regarding eligibility criteria, dependant coverage, claims process and more (from the uploaded knowledge).

2. **Enable collaboration between AskBenefits and pre-configured AskDental agent** – an agent that supports questions about dental benefits by referencing a PDF document. It can answer a range of questions about coverage and plan details.

What you'll need

- The OpenAPI specification file (ask_benefits_tools.yaml) – Click here to download
- The HR policy document on medical benefits (HR_Policy_Medicalbenefits.pdf) – Click here to download
- The dental benefits file (dental_benefits_summary.pdf) – Click here to download

Part 1: Build and engage with a new AskBenefits agent

Step 1. Create a new AskBenefits agent

- Click on '**Create new Agent'** in the bottom left corner

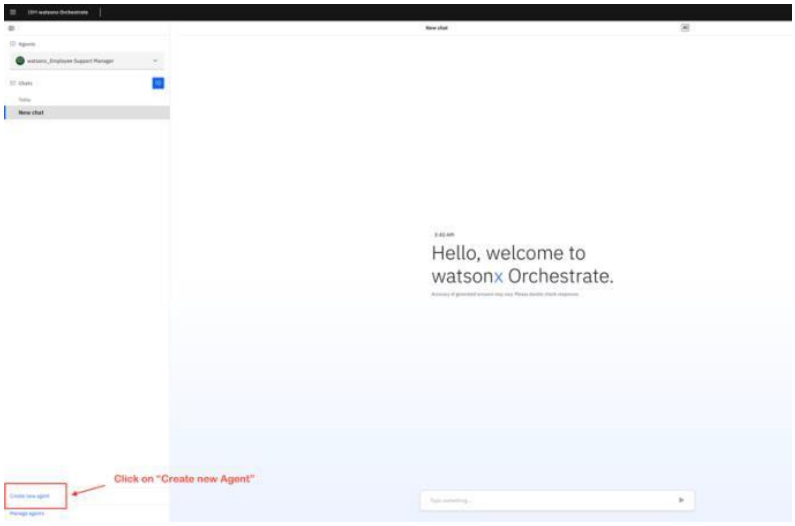

- Enter the Agent Name and description:

- Agent Name – <Your_Name>_AskBenefits
- Agent Description – A smart healthcare agent that classifies medical queries and uses the right tools to retrieve data, guide care, book appointments and explain coverage with clarity and empathy.

- Then click on '**Create**'.

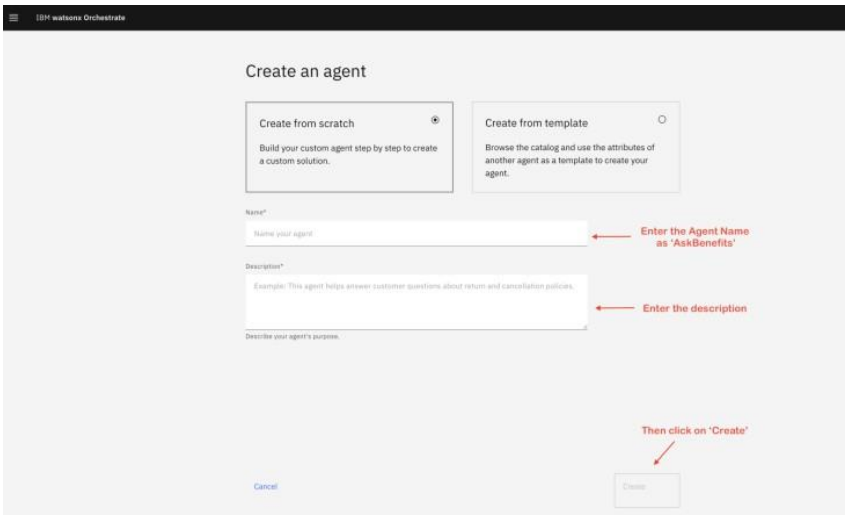

Step 2. Import and attach the tools

- Click the '**Toolset**' tab on the left-hand panel.

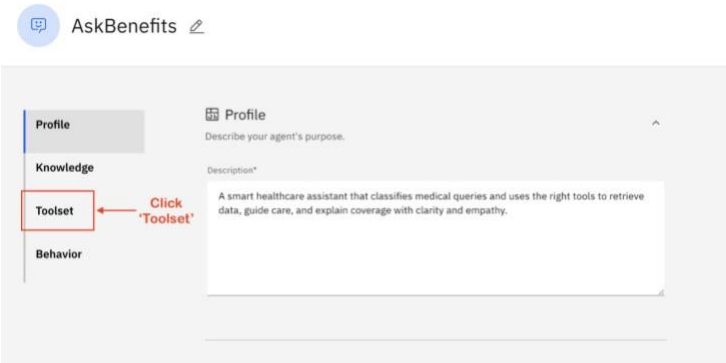

- Click '**Add tool**' button.

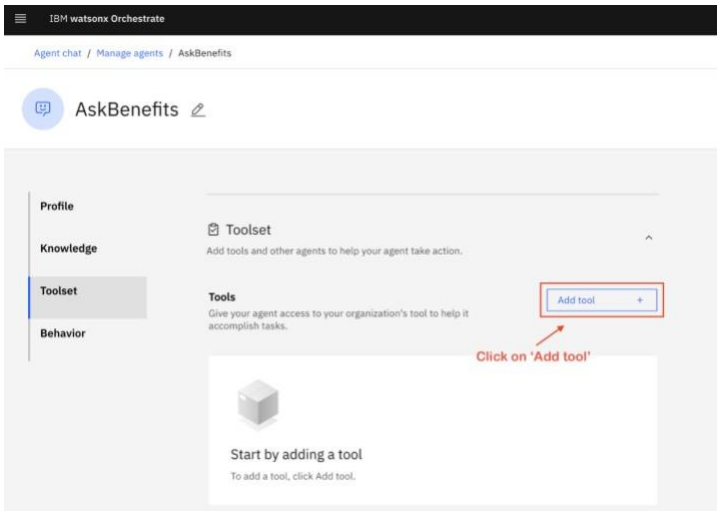

- Select the '**Import**' option in the pop-up window.

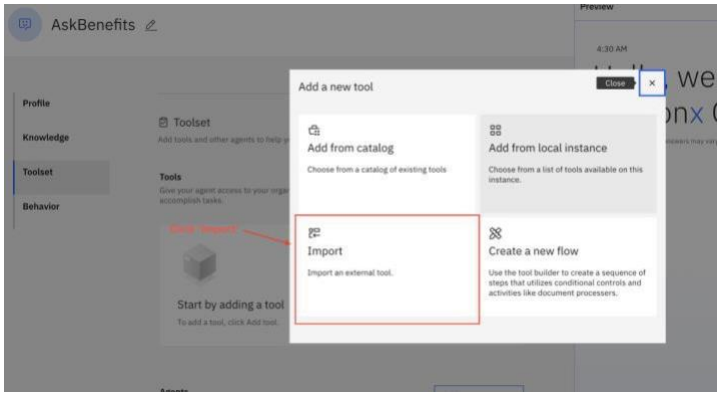

- Select '**Import from file**' option.

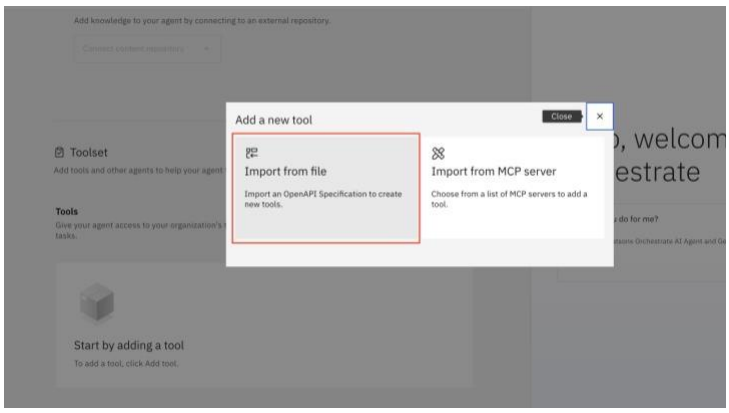

- Click inside the blue box to select the '**ask_benefits_tools.yaml**' file downloaded from the **What you'll need** section above.

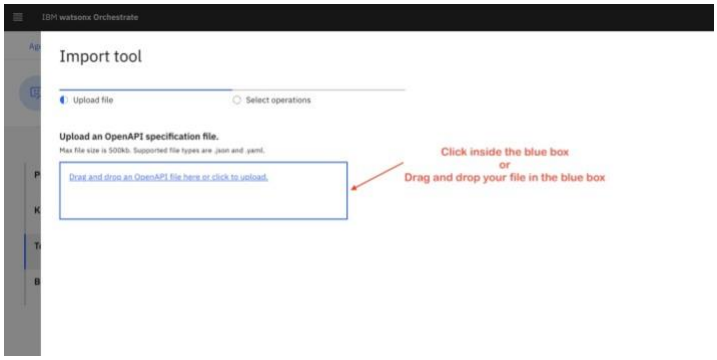

- Once the file is imported successfully, click '**Next**'.

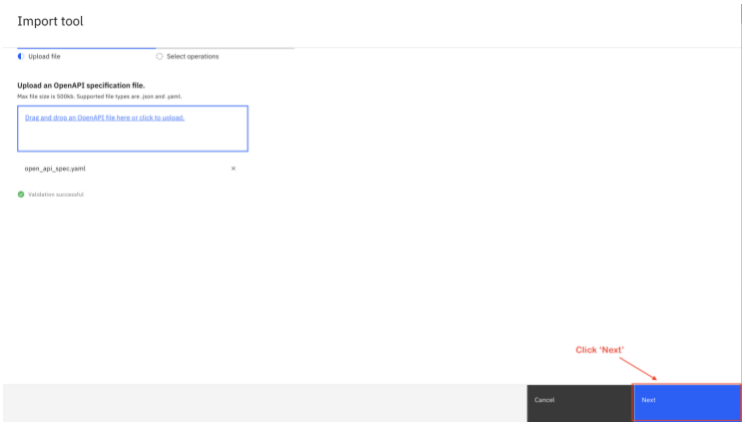

- Click '**Select all rows in the table**' checkbox to select all available tools.

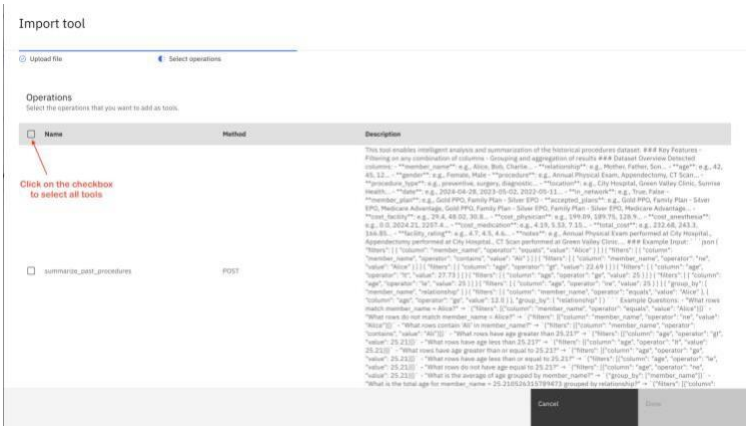

- Then click on the blue '**Done**' button at the right bottom corner.

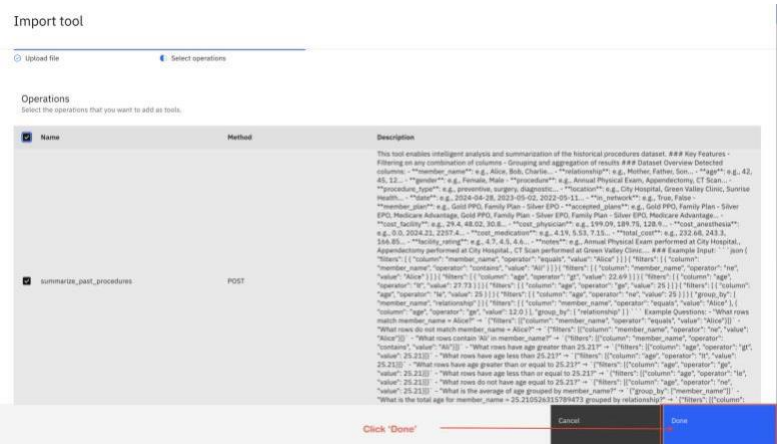

- The 9 tools have now been added to your Agent.

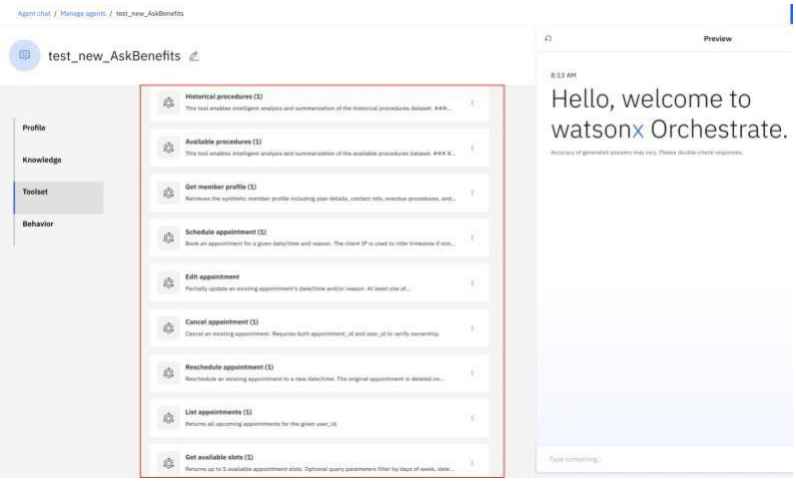

Step 3. Upload knowledge document

- Click the '**Knowledge**' tab in the left-hand panel.

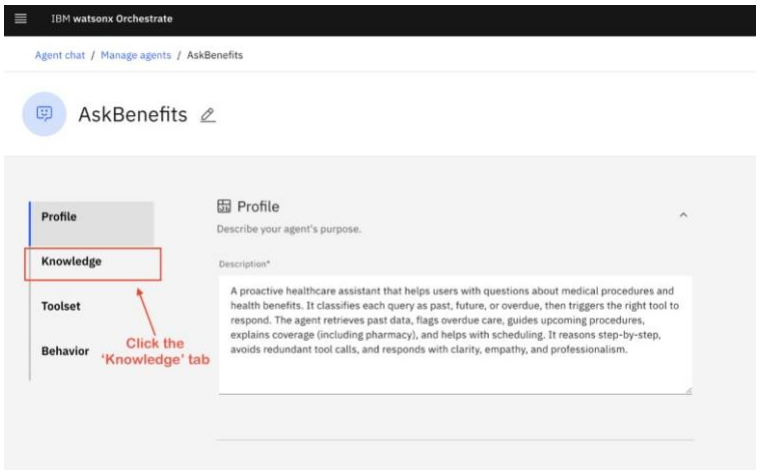

- In the '**Document**' section click on '**Upload files**'.

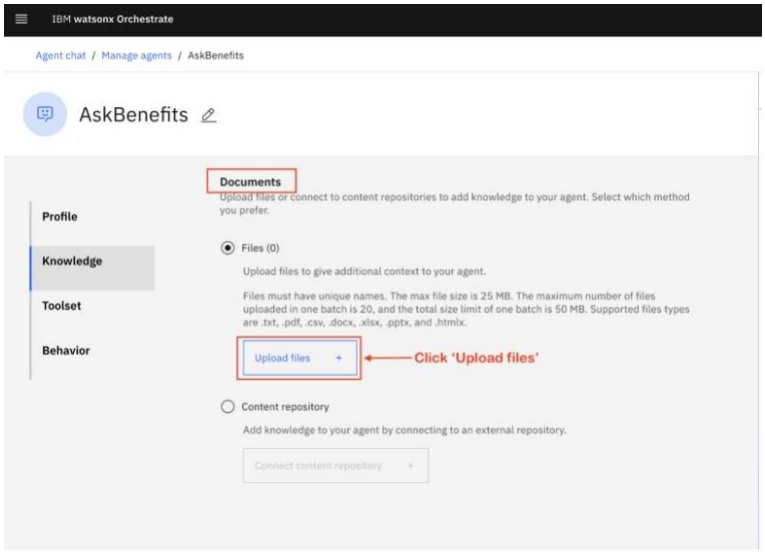

- Select the **HR_Policy_Medicalbenefits.pdf** file downloaded from the '**What you'll need**' section.

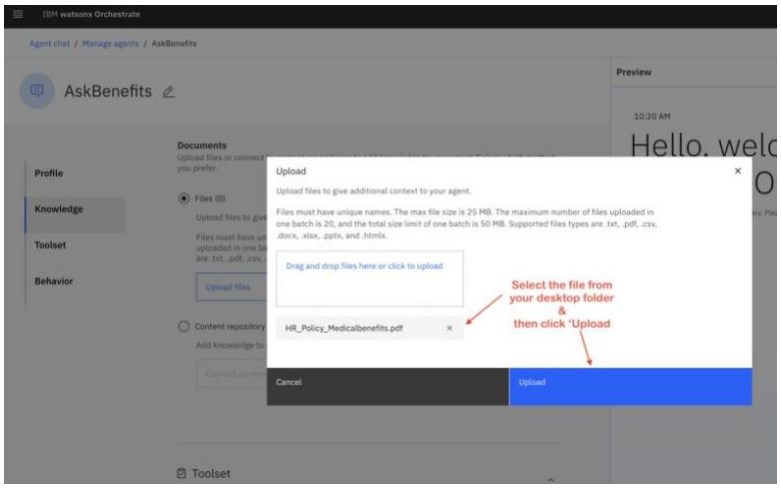

- As a best practice, in the '**Knowledge**' description field, add the below context to inform the agent on the contents of the knowledge:

  "*This document covers key healthcare benefit policies—eligibility, coverage, claims, exclusions, and post-employment options. Use it to answer employee questions related to health benefits*"

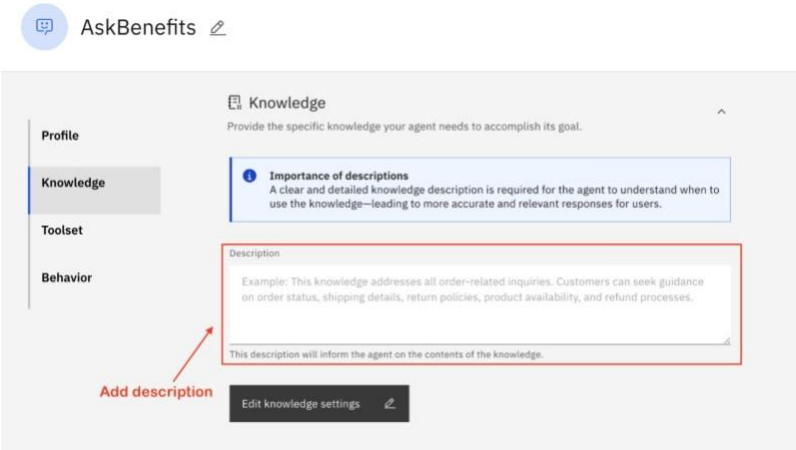

Step 4. Test the agent

In the Preview panel enter the below prompts one-by-one in the "Type something..." field and press **Enter**.

- **Prompt 1**: Can you give me a cost breakdown for X rays?

  ***Expected outcome**:* The agent displays the X-ray cost breakdown across multiple hospitals, including facility, physician, anesthesia, and medication charges.

  What's happening behind the scenes? To understand that click '**Show reasoning**' and expand the steps.

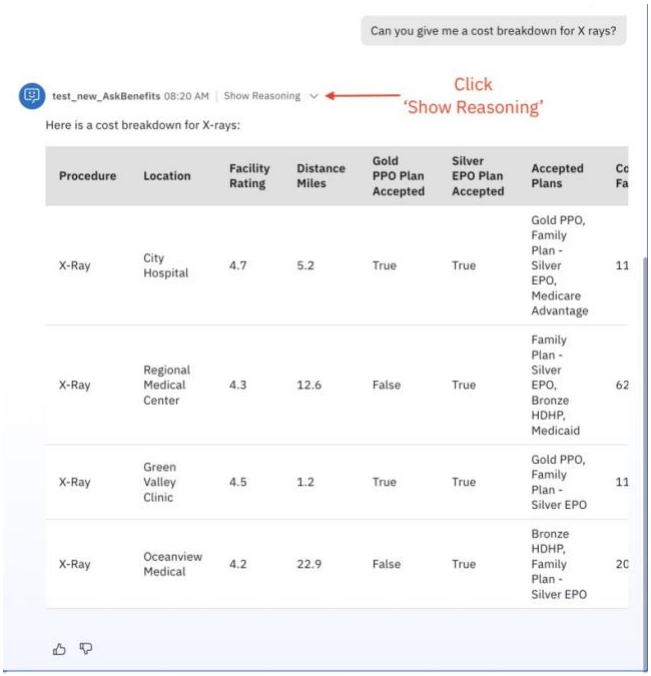

  ***Key Observations**:* When you insert the prompt, the agent:

  1. Understands User Intent: The agent correctly interprets the request as requiring a filtered view of cost-related procedure data.

1. Dynamic Tool Invocation: It selects the correct tool '*Available_procedures'*, which is appropriate for analyzing procedure costs from a dataset.

2. StructuredReasoningviaFilters:Itcreatesastructuredqueryusing:

```
"filters": "[{\"column\": \"procedure\", \"operator\": \"equals\", \"value\": \"X ray\"}]"
```

This shows the agent's ability to construct a parameterized request dynamically based on the user's natural language.

- **Prompt 2:** *Which among them is the cheapest?*

  Expected Outcome: The agent compares total costs and responds with the hospital offering the lowest price.

  Key Observation: This is a strong example of co-referencing, where the agent correctly interprets the user's use of "them" based on prior context. It also demonstrates the use of memory to retain and build on prior context—both key aspects of agentic behaviour.

- **Prompt 3**: Am I eligible to add my spouse to my health insurance?

  Expected Outcome*:* The agent confirms that spouses are eligible dependents and advises the employee to provide supporting documents during enrolment.

  Key Observation*:* The agent intelligently processes the document, identifies the relevant section on dependent coverage, and retrieves a precise answer—demonstrating context-aware document parsing and retrieval.

Step 5. Testing the agent

- Click "**Deploy**" to publish your AskBenefits agent.

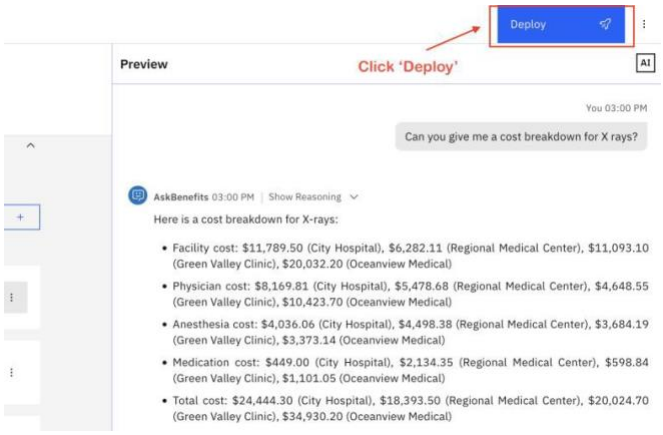

- Congratulations! You successfully deployed the AskBenefits Agent.

Step 6. Identify the knowledge gap

- Click on '**Chat**' in the Navigation menu.

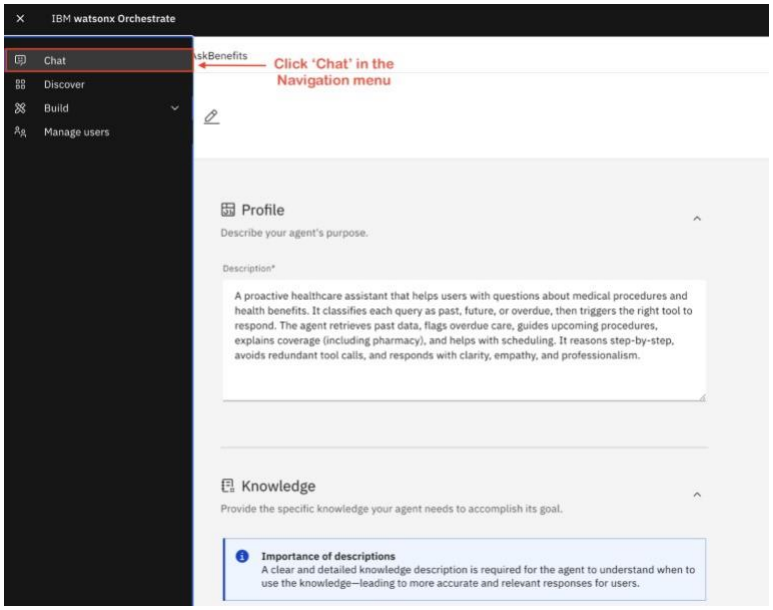

- Select the '**AskBenefits**' agent in the Agent dropdown and enter the below prompt in the 'Type something…' field. Then press '**Enter**'.

  Prompt: *Can you describe my Dental benefits?*

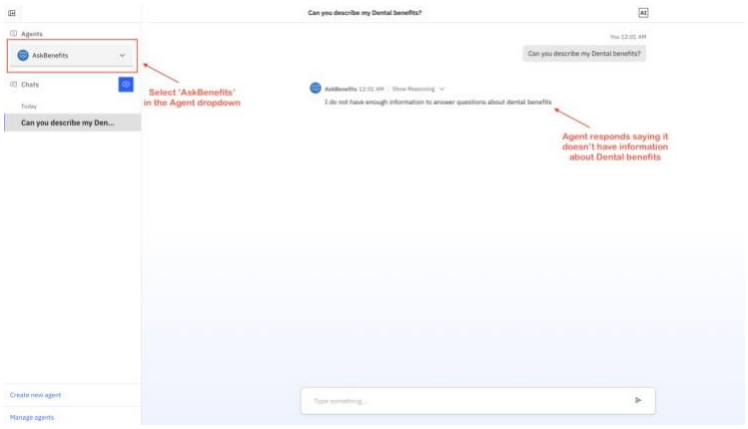

  Agent response: *I do not have enough information to answer questions about dental benefits.*

- Let's now look at how the above information can be retrieved from the AskBenefits agent. A collaborator agent is used by AskBenefits to retrieve answers related to Dental benefits.

Part 2: Create a new collaborator agent – AskDental

Step 1. Create a new AskDental agent

- Click on '**Create new Agent**' in bottom left corner.

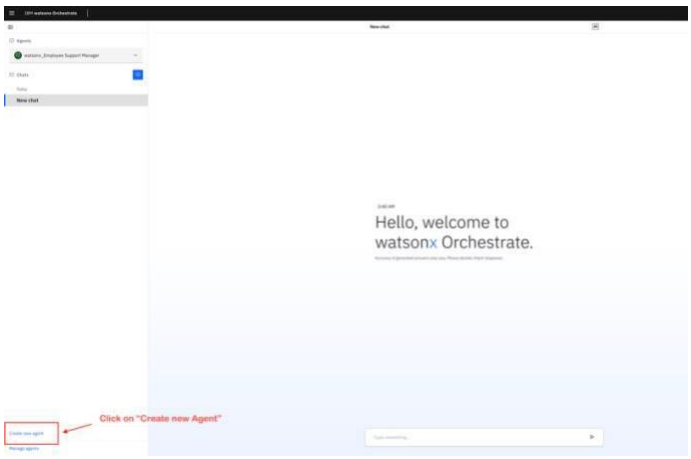

- Enter the Agent Name and description:
  1. Name – AskDental
  2. Agent Description – The dental benefits agent is an AI assistant that helps policyholders understand their coverage and make informed oral health decisions. It analyzes individual policy details to answer questions about dental benefits.

- Then click on '**Create**'.

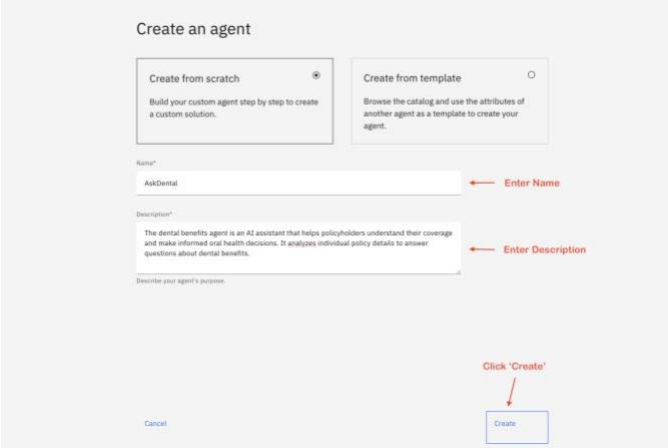

- Click '**Upload files**' under the '**Knowledge**' tab

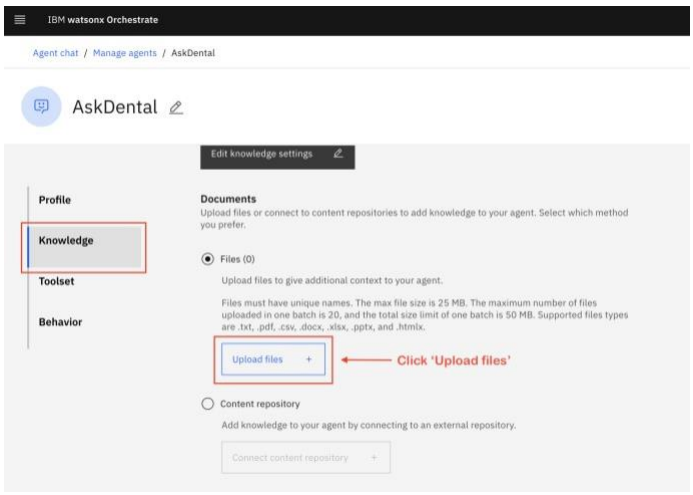

- Select the file containing **Dental Benefit information *dental_benefits_summary.pdf*** you downloaded from What *you'll need* section.

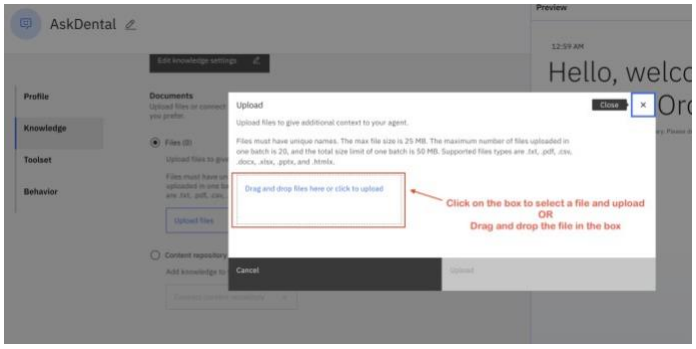

- Click '**Upload**'. The document upload process can take a few seconds.

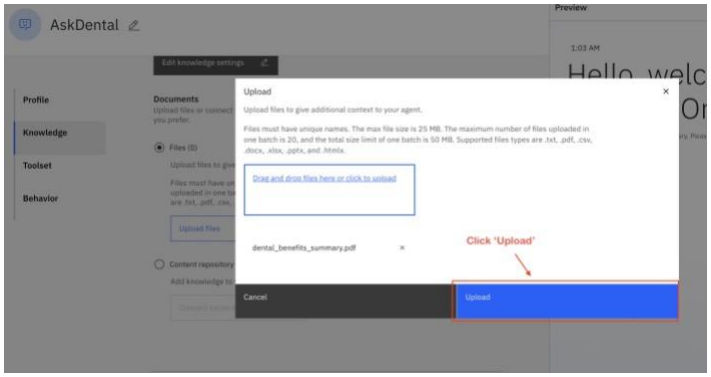

- In the Knowledge '**description**' field enter the following description:

  "This knowledge file outlines dental benefits and is used to answer policyholder questions. It includes details about individual coverage and policies. If a question cannot be answered using the available content, do not respond directly. Instead, forward the query verbatim to the supervisor agent for further handling."

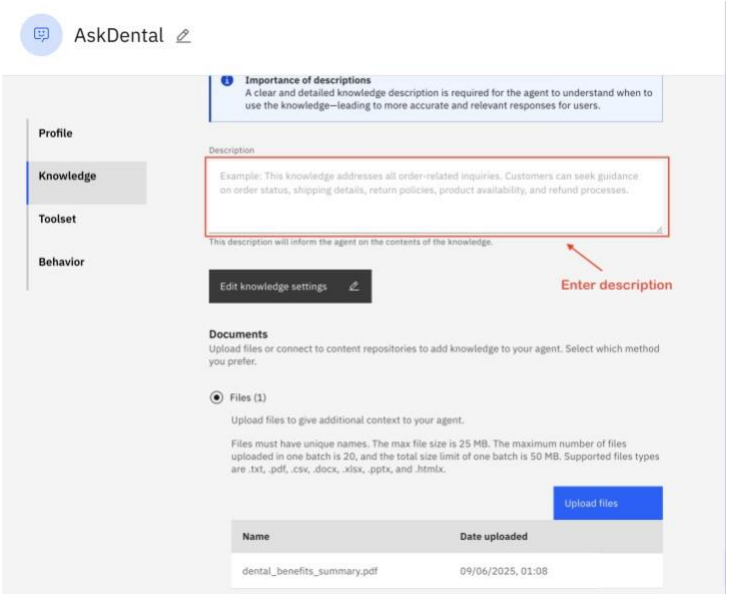

- Click '**Deploy**'.

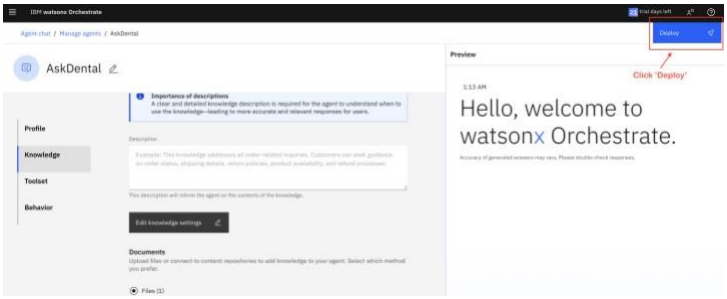

Step 2. **Add AskDental agent as a collaborator**

- Click '**Manage agents**'.

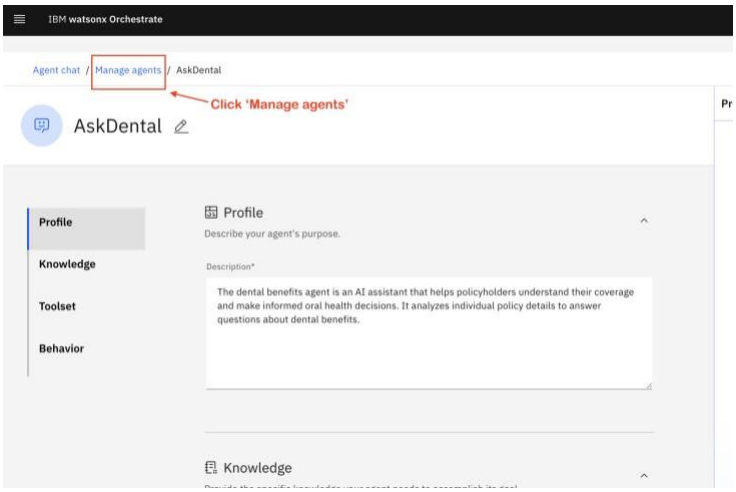

- Select '**AskBenefits**'.

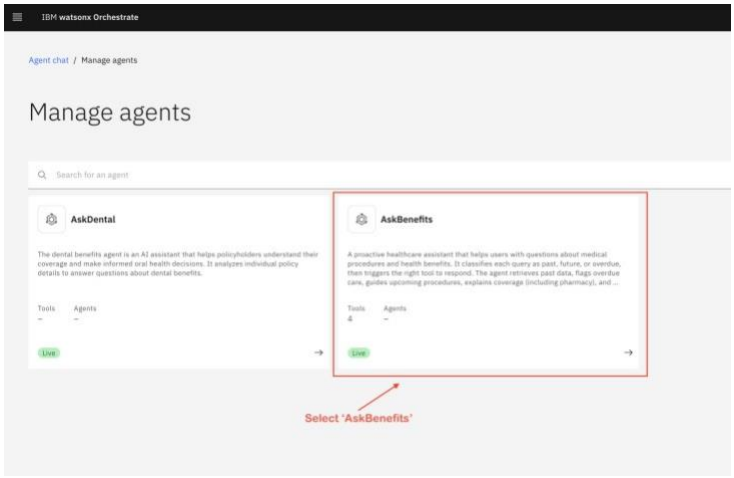

- Click '**Add agent**' under the '**Toolset**' tab.

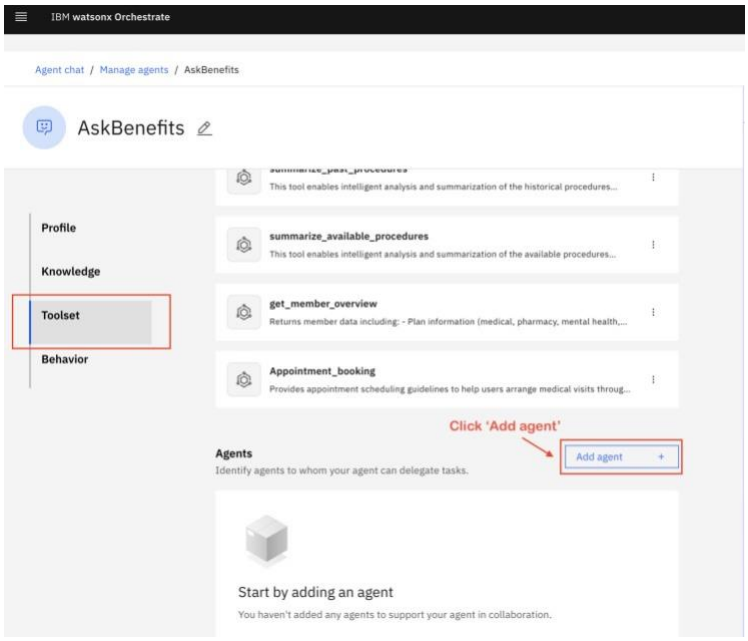

- Click '**Add from local instance**'.

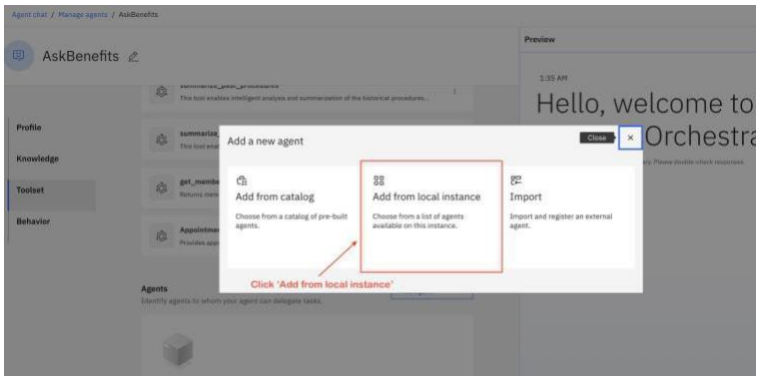

- Select '**AskDental**' agent and then click '**Add to agent**'.

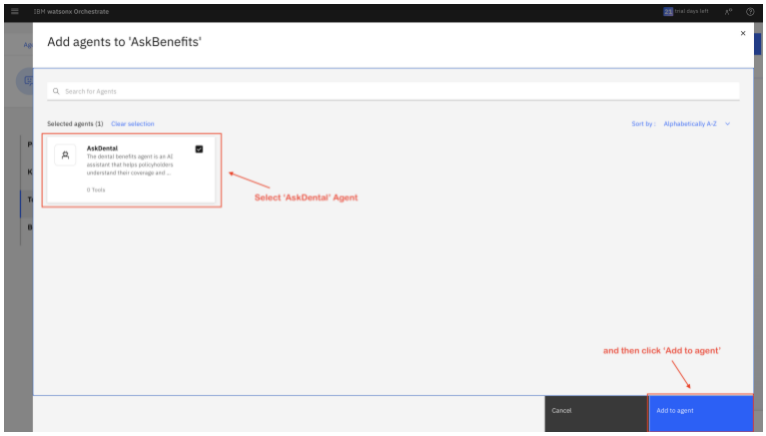

- Click '**Deploy**'.

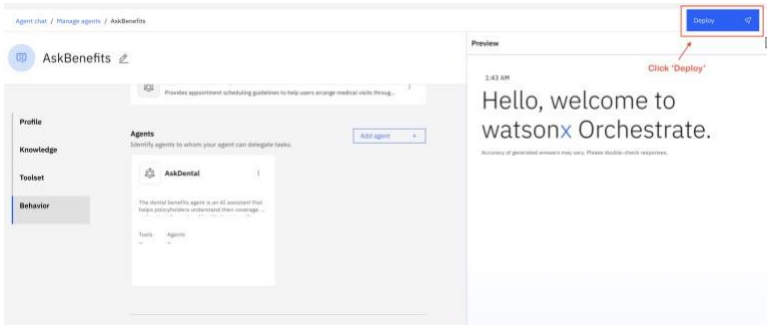

- Click '**Chat**' to navigate to the Agent chat screen.

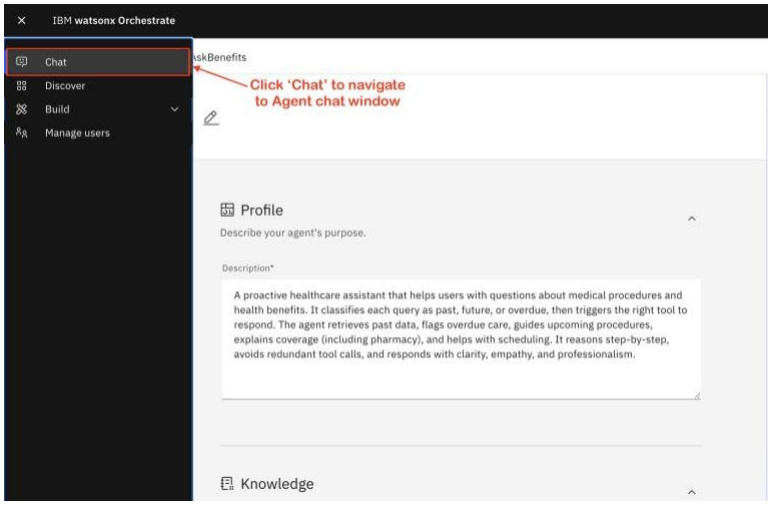

Step 3. **Testing scenarios**

After successful deployment, select the '**AskBenefits**' agent in the Agent dropdown and test the same prompt to verify how the collaborator agent added to AskBenefits helps in answering Dental related queries.

- **Prompt**: *Can you describe my dental benefits?*

  ***Expected Outcome**:* The AskBenefits agent routes the query to the AskDental agent, which

responds with a detailed summary of dental benefits including deductible, coverage, and plan-specific details—retrieved from the uploaded dental policy document.

***Key Observation**:* The AskBenefits agent correctly routes the dental benefits query to the AskDental agent, which retrieves a detailed response from its underlying document—demonstrating effective agent collaboration and document-grounded reasoning.

- **Try some more prompts:**
    1. What procedures have we had at City Hospital in the last year?
    2. Show some appointment slots available for Dental checkup?
    3. Can you show the available slots on weekends instead?
    4. Book an appointment for <Enter an available slot>
    5. List my upcoming appointments
    6. Give me a summary of my insurance plan